

Comparing Interpretable Sequence Features and Protein Language Models for Missense Variant Classification

Georgios Ioannou (gi2100)

Vedant Jagtap (vsj7589)

AI in Genomics, Spring 2026

Background and motivation

A *genetic variant* is a place in a person’s DNA that differs from the human reference. We focus on *missense variants*: DNA changes that swap one amino acid for another in the protein; some are disease-causing (*pathogenic*) and others are tolerated (*benign*). **ClinVar** [2, 6] is our primary label source. **ESM-2** [3, 9] is a protein language model whose *residue embeddings* (1280-d for the 650M checkpoint) summarize evolutionary context. We ask, on a **gene-held-out** split, (i) how much signal lives in small interpretable feature sets and (ii) whether frozen embeddings improve cross-gene AUROC/AUPRC beyond those features—the tradeoff among interpretability, compute cost, performance, and generalization.

Methods

Data. We use NCBI’s ClinVar `variant_summary.txt` (release dated **2026-05-04**, the snapshot pinned in `data/processed/clinvar_release.txt`) [5]. We restrict to GRCh38 germline single-nucleotide missense variants whose ClinVar review status maps to a clear binary label: pathogenic / likely-pathogenic vs. benign / likely-benign [4, 7]. We exclude conflicting and uncertain interpretations. Duplicates (same RefSeq transcript, position, ref AA, alt AA) are dropped, keeping the strongest review status. Table 4 reports the per-stage retention counts. We map each retained variant to a *reviewed* (Swiss-Prot) human protein sequence in UniProt UP000005640 [10], requiring that the residue at the variant position in the canonical sequence equals the HGVS-reported reference amino acid. After this verification we retain **193,483** variants across **14,893** genes (62,241 pathogenic / 131,242 benign). Figure 5 sketches the full pipeline.

Splits. To prevent *train–test contamination across variants of the same gene*, we use a **gene-disjoint** 70/10/20 train/val/test split (seed 42), stratified across genes by majority label so neither split is all-positive or all-negative. Final split sizes: **train 135,418** / **val 19,539** / **test 38,526** variants, with 10,425 / 1,489 / 2,979 unique genes respectively.

Models. We train five predictors: two baselines on handcrafted features, two neural heads on ESM-2 embeddings (embedding-only and concatenated with handcrafted features), plus AlphaMissense [1] and CADD [8] as external references (not targets to outperform). The handcrafted vector per variant is exactly the five proposal families (20-d reference AA one-hot, 20-d alternate one-hot, BLOSUM62 score, 20-d local composition with window radius 7, normalized position; 62-d total). Table 1 summarizes inputs and architectures; tile handling for long proteins and optimization hyperparameters are deferred to Appendix .

Table 1: **Our models at a glance.** Neural heads keep ESM-2 650M frozen; training uses 135,418 gene-disjoint variants.

Model	Input	Architecture	Training
Logistic re- gression	62	L2 LR, balanced, C tuned on val AUPRC	Train split only
Random for- est	62	500 trees, balanced, depth/leaf grid on val AUPRC	Train split only
ESM-2 head	1280	MLP $d \rightarrow 256 \rightarrow 1$, Layer-Norm, dropout 0.2, GELU	AdamW + cosine (Appendix)
Combined head	1342	Same MLP with $d = 1342$, concat(handcrafted, emb.)	Same as ESM-2 head

ESM-2 caveat (proposal-flagged). ESM-2 was pretrained on UR50, which contains many human proteins; this could inflate metrics on genes already well-represented in pretraining. We surface this caveat here and analyze it empirically in the appendix (Figure 8). **Computational simplicity** is reported per model in Appendix Table 5.

Evaluation. Metrics mirror Phase 9: AUROC, AUPRC, F1 at the val-tuned threshold, and 95% bootstrap CIs (1000 resamples). AlphaMissense and CADD are scored on the Phase 10 intersections with the test variants.

Results

Both empirical questions answered by the numbers in Table 2. The interpretable handcrafted random forest reaches test AUROC = 0.850 and AUPRC = 0.764 – roughly *nine-tenths of the ESM-2 head’s AUROC* and *85% of its AUPRC* on a strict cross-gene test. A paired bootstrap on identical variants confirms the ESM-2 head’s AUROC advantage over that forest ($\Delta = +0.089$, 95% CI [0.085, 0.093], two-sided $p < 0.001$). Concatenating handcrafted features with embeddings yields a further increment (combined head: AUROC 0.954, +0.015 vs. embedding-only), indicating partial complementarity. AlphaMissense remains the strongest reference (AUROC 0.966); the ESM-2 head and CADD bracket it from below (Figure 1).

Gene-level and error structure. Per-gene AUROCs and confusion summaries (Figure 2) show the ESM-2 head’s gene-level median above both baselines with a tighter spread, balanced precision/recall (0.80 / 0.84) versus LR (0.55 / 0.78) and RF (0.64 / 0.75). The combined head reaches the highest F1 among our models (0.85 at the val-tuned threshold).

Discussion

Direct numerical answers to the proposal’s two questions. (i) *How much pathogenicity signal do simple interpretable features capture?* The L2 logistic regression reaches AUROC 0.80, and the random forest on the five headline families reaches AUROC 0.85 / AUPRC 0.76. A feature-

Table 2: **Test-set metrics with 95% bootstrap CIs.** Main test $n = 38,526$; AlphaMissense $n = 37,736$; CADD $n = 38,792$ (Phase 10 matching).

Model	n	AUROC [95% CI]	AUPRC [95% CI]	F1 [95% CI]	Acc.
Logistic regression	38,526	0.800 [0.796, 0.805]	0.668 [0.659, 0.677]	0.647 [0.641, 0.653]	0.708
Random forest	38,526	0.850 [0.846, 0.854]	0.764 [0.757, 0.771]	0.691 [0.684, 0.697]	0.768
ESM-2 head	38,526	0.939 [0.936, 0.941]	0.901 [0.897, 0.906]	0.820 [0.815, 0.825]	0.874
Combined head	38,526	0.954 [0.952, 0.956]	0.924 [0.921, 0.928]	0.850 [0.845, 0.854]	0.896
AlphaMissense (ref.)	37,736	0.966 [0.964, 0.967]	0.939 [0.935, 0.942]	0.871 [0.867, 0.876]	0.912
CADD (ref.)	38,792	0.928 [0.925, 0.931]	0.839 [0.832, 0.846]	0.807 [0.803, 0.812]	0.860




Figure 1: **Test-set ROC and PR curves (Figure 2).** All models on $n = 38,526$ gene-held-out test variants (AlphaMissense on 37,736 with score coverage). The ESM-2 head sits between the handcrafted baselines and AlphaMissense.

family ablation at the Phase 7 RF hyperparameters (Appendix Figure 4) shows that reference + alternate one-hot encodings plus the scalar BLOSUM62 term alone achieve test AUPRC 0.622, i.e. $\sim 81\%$ of the full forest’s AUPRC (0.764), with almost all remaining lift coming from the 20-d local composition window. (ii) *Does ESM-2 improve cross-gene generalization beyond handcrafted features?* Yes: the embedding-only head gains AUROC +0.089 and AUPRC +0.138 over the RF on mutually exclusive genes, and the paired bootstrap on identical variants confirms the AUROC contrast ($\Delta = +0.089$, 95% CI [0.085, 0.093], two-sided $p < 0.001$). Concatenating embeddings with handcrafted vectors nudges performance upward again (Table 2), consistent with partially complementary information.

Tradeoffs, calibration, and limitations. Interpretability and disk footprint favor logistic regression; the forest remains expensive at inference; neural heads require a one-time embedding pass but rank best among our models. Reliability diagrams (Appendix Figure 3) show that all scorers are reasonably well centered, with the neural heads achieving the lowest Brier scores (Table 3). Limitations include canonical isoform mapping (non-canonical transcripts are dropped), UR50 pretraining overlap with human proteins, and overlap between external tools and ClinVar—AlphaMissense and CADD should be read as optimistic references. F1 uses val-tuned thresholds; the headline story is threshold-free AUROC/AUPRC.

Future work. End-to-end ESM-2 fine-tuning on ClinVar labels or structured clinical loss functions could probe how much of the current gap to AlphaMissense is fixable with data rather than




Figure 2: **Test-set confusion matrices (top) and per-gene AUROC distribution (bottom).** Top: val-tuned thresholds. Bottom: 247 genes with ≥ 5 pathogenic and ≥ 5 benign variants.

architecture.

Supplementary materials

supplementary_team19_gi2100_vs_j7589.zip bundles code, configs, pinned dependencies, REPRODUCE.md, data README, all tables/figures (including pairwise_tests.csv, calibration, ablations), and the GitHub mirror <https://github.com/GeorgiosIoannouCoder/interpretable-missense-classification>. Run `make reproduce`; checkpoints stay resumable.

References

- [1] Jun Cheng, Guido Novati, James Pan, Clare Bycroft, Akvilė Žemgulytė, Taylor Applebaum, Alexander Pritzl, Lily H. Wong, Michal Zielinski, Toby Sargeant, Richard G. Schneider, Andrew W. Senior, John Jumper, Demis Hassabis, Pushmeet Kohli, and Žiga Avsec. Accurate

- proteome-wide missense variant effect prediction with AlphaMissense. *Science*, 381:eadg7492, 2023. doi: 10.1126/science.adg7492.
- [2] Melissa J. Landrum, Shanmuga Chitipiralla, Gary R. Brown, Chao Chen, Baoshan Gu, Jennifer Hart, Douglas Hoffman, Wonhee Jang, Kuljeet Kaur, Chunlei Liu, Vitaly Lyoshin, Zenith Maddipatla, Rama Maiti, Joseph Mitchell, Nuala O’Leary, George R. Riley, Wenyao Shi, George Zhou, Valerie Schneider, Donna Maglott, J. Bradley Holmes, and Brandi L. Kattman. ClinVar: improvements to accessing data. *Nucleic Acids Research*, 48(D1):D835–D844, 2020. doi: 10.1093/nar/gkz972. URL <https://doi.org/10.1093/nar/gkz972>.
 - [3] Zeming Lin, Halil Akin, Roshan Rao, Brian Hie, Zhongkai Zhu, Wenting Lu, Nikita Smetanin, Robert Verkuil, Ori Kabeli, Yaniv Shmueli, Allan dos Santos Costa, Maryam Fazel-Zarandi, Tom Sercu, Salvatore Candido, and Alexander Rives. Evolutionary-scale prediction of atomic-level protein structure. *Science*, 379(6637):1123–1130, 2023. doi: 10.1126/science.ade2574.
 - [4] NCBI ClinVar. Representation of classifications in ClinVar. <https://www.ncbi.nlm.nih.gov/clinvar/docs/clinsig/>, 2026. Official classification terminology page.
 - [5] NCBI ClinVar. Guide to using files from the FTP site or accessed via e-utilities. https://www.ncbi.nlm.nih.gov/clinvar/docs/ftp_primer/, 2026. Official description of variant_summary.txt.
 - [6] NCBI ClinVar. What is ClinVar? <https://www.ncbi.nlm.nih.gov/clinvar/intro/>, 2026. Official overview page.
 - [7] NCBI ClinVar. Review status in ClinVar. https://www.ncbi.nlm.nih.gov/clinvar/docs/review_status/, 2026. Official review-status documentation.
 - [8] Philipp Rentzsch, Daniela Witten, Gregory M. Cooper, Jay Shendure, and Martin Kircher. CADD: predicting the deleteriousness of variants throughout the human genome. *Nucleic Acids Research*, 47(D1):D886–D894, 2019. doi: 10.1093/nar/gky1016.
 - [9] Alexander Rives, Joshua Meier, Tom Sercu, Siddharth Goyal, Zeming Lin, Jason Liu, Demi Guo, Myle Ott, C. Lawrence Zitnick, Jerry Ma, and Rob Fergus. Biological structure and function emerge from scaling unsupervised learning to 250 million protein sequences. *Proceedings of the National Academy of Sciences*, 118(15):e2016239118, 2021. doi: 10.1073/pnas.2016239118. URL <https://www.pnas.org/doi/10.1073/pnas.2016239118>.
 - [10] UniProt. Homo sapiens reference proteome (UP000005640). <https://www.uniprot.org/proteomes/UP000005640>, 2026. Official proteome page.

Appendix

Neural training details

Both the ESM-2 and combined heads are trained with AdamW (learning rate 10^{-3} , weight decay 0.01), batch size 512, BCE-with-logits with a positive-class weight $N_{\text{neg}}/N_{\text{pos}}$ estimated on the training split, and a cosine learning-rate schedule spanning all optimizer steps. Mixed precision (bf16) is enabled on CUDA when available via Hugging Face Accelerate (so the same code spans multi-GPU servers). Early stopping tracks validation AUPRC with patience five epochs; checkpoints

are written every 500 steps and at epoch boundaries. For proteins exceeding ESM-2’s 1022-token context we slide overlapping sequence windows, keep all variant-containing tiles, and select the tile whose center lies closest to the missense site before pooling the residue hidden state.

Table 3: **Table A2 / Calibration on the gene-held-out test set** ($n = 38,526$). Brier score and expected calibration error (ECE) with ten quantile bins; lower is better. Plots appear in Figure 3.

Model	Brier	ECE
Logistic regression	0.182	0.113
Random forest	0.148	0.040
ESM-2 head	0.093	0.049
Combined head	0.080	0.048

Table 4: **Table 1 / Filtering funnel**. Each row reports retained ClinVar variants after applying the labelled stage. The proposal’s cited 2026-03-30 baseline (4.24 M variation records) is reproduced approximately at the GRCh38 stage with our 2026-05-04 snapshot.

Stage	Variants retained
Raw <code>variant_summary.txt</code> rows	8,978,110
After Assembly = GRCh38	4,455,427
After germline only	4,142,993
After single-nucleotide variant type	3,859,832
After missense parse	2,415,038
After P/LP or B/LB label only	210,038
After dedup by (transcript, AA pos, ref AA, alt AA)	208,490
After UniProt Swiss-Prot mapping with ref-AA verification	193,483

Table 5: **Table A1 / Computational-simplicity (efficiency) per model**. Inference latency measured per variant on the GH200 (CPU for sklearn models, GPU for the ESM-2 head – excluding the one-time 4 minutes of embedding extraction).

Model	Params	Disk size (MB)	Train time (s)	Inference (ms / variant)
Logistic regression	63	0.001	0.4	0.04
Random forest	500 trees	2,412.7	1.8	44.3
ESM-2 head	331,265	3.8	16.5	0.22
Combined head	347,261	4.0	11.4	0.20




Figure 3: **Figure A2 / Reliability diagrams.** Calibration curves for the four in-house scorers (10 quantile bins). Points should hug the diagonal under perfect calibration.




Figure 4: **Figure A3 / Random forest feature-family ablation.** Test AUPRC when training only on cumulative feature subsets (top block) or removing one family from the full 62-d vector (bottom block) using the Phase 7 hyperparameters. The local window composition drives most of the lift beyond substitution-aware features.




Figure 5: **Figure 1 / Pipeline schematic.** Data flow from raw ClinVar to the gene-held-out evaluation.




Figure 6: **Figure 4 / Handcrafted-feature interpretability.** Top-15 LR coefficients (left) and top-15 RF Gini importances (right). The most influential features are alternate amino acids (Cys, Gly, Pro, Trp), the BLOSUM62 score, and one-hot reference amino acids – all biologically expected.




Figure 7: **Figure 5 / UMAP of ESM-2 test embeddings.** A 5,000-variant random subsample of the test set, colored by ClinVar label. Pathogenic and benign clouds are partly separable in the embedding space even before any classifier is trained – the head’s job is largely a linear cleanup.




Figure 8: **Figure 7 / Per-gene AUROC vs gene studiedness.** A linear fit of per-gene AUROC versus $\log(1 + \text{ClinVar variants per gene})$ (a proxy for how heavily a gene is studied). The ESM-2 head's slope is not noticeably steeper than the baselines', suggesting its advantage is broadly distributed rather than concentrated on well-represented genes.




Figure 9: **Figure 8 / Review-status sensitivity.** Test AUROC and AUPRC stratified by ClinVar review-status confidence: full retained set, ≥ 2 -star (criteria, multiple submitters, no conflicts), and expert (≥ 3 -star: reviewed by expert panel or practice guideline). Conclusions hold across strata; the small expert-only stratum is more variable.